

MPLADS-SATARK

Final Solution Blueprint — Real-Data Edition

Problem Statement 26102 · Ministry of Statistics & Programme Implementation (MoSPI / DIID)
Development of an AI-powered system to detect anomalies, fraud, and inefficiencies in MPLAD Scheme implementation

VERSION 5 — NATIONAL ROUND EDITION

What Changed in v5

1. **Feasibility, engineered away.** A bifurcated cold-path / hot-path compute architecture (DuckDB + pre-fitted model scoring + ego-graph rendering + probabilistic record linkage) removes every scalability risk that could break a live demo of 73,000+ records.
2. **Sustainability, engineered in.** Every statutory rule is now Policy-as-Data with bitemporal versioning — the system survives future changes to GFR 2017 / MPLADS Guidelines without a single code redeploy, and every flag stays audit-defensible years later.
3. **A fifth independent fraud signal.** Benford's Law forensic digit analysis — the same class of technique used in real tax and audit investigations — added as a statistically independent layer alongside the ML and graph signals.
4. **Real prediction, not a trend line.** Survival analysis (Cox Proportional Hazards) now forecasts project completion probability and flags "zombie project" risk months earlier than a naive delay counter.

1 The Problem, In Plain Words

Every year, Members of Parliament receive funds — called MPLADS funds, over Rs. 4,000 Crore across 788 districts — to build things people need in their area: roads, community halls, drinking water systems, school repairs, and more. Right now, the official government dashboard only shows what already happened. It doesn't warn anyone before money gets wasted on duplicate projects, overpriced work, or long delays.

MPLADS-SATARK is our answer: a system that reads real government data, spots unusual patterns early, enforces statutory financial rules, and gives officers — and every level of authority above them — a clear, explainable list of what to check first, before public money is fully spent.

2 What Makes Us Different

- ✓ **Real Data Only** — we only use real government data. No synthetic numbers, no invented contractor IDs, no fabricated GPS points.
- ✓ **Honest Labeling** — every number is labeled Real, Derived, or Estimated, so nobody, including a judge, is ever misled.
- ✓ **Statutory Alignment** — grounded in GFR 2017 and the MPLADS 2023 Guidelines, with every citation checked against the primary source.
- ✓ **Human In The Loop** — our AI never has the final say. It only recommends. A human officer always makes the administrative decision.
- ✓ **Zero Black Boxes** — every score carries an itemized, additive "Explain-My-Score" breakdown traceable to ground evidence.
- ✓ **Engineered For Failure Modes** — the architecture is explicitly designed against the three things that break real government deployments: scale, changing rules, and cross-ministry data quality. Sections 7 and 8 show exactly how.

3 The Real Data We Use

Data Source	What It Tells Us	Where It Comes From
MPLADS Work Records (73,000+ real projects)	Project name, cost sanctioned, MP, district, dates, current status, sector	Official MoSPI records (mplads.gov.in)
MPLADS Fund-Flow Data	How much money each MP received, released, and left unspent — year by year	Official MoSPI records
MGNREGA / PMGSY Asset Geotags	Real GPS locations of other government-built assets nearby, for duplicate cross-checking	Government geo-tagging system (Bhuvan / ISRO)
CPWD Cost Rate Book (DSR)	Official standard construction costs to benchmark and spot overpriced projects	Central Public Works Department (CPWD)

Statutory Rules & Guidelines	Financial rules, negative list of ineligible works, and past audit findings, each cross-checked against its primary source	GFR 2017 & MPLADS Guidelines 2023 (mplads.gov.in)
------------------------------	--	---

4 Data Provenance Framework

We label every field so nothing is ever mistaken for something it isn't — an auditor's first question is always "where did this number come from?"

Label	Meaning
REAL	Taken directly from official government records. Nothing calculated or imputed.
DERIVED	Calculated from real data using a clear, explainable formula, shown to officers on request.
ESTIMATED	A transparent, rule-based approximation, used only when exact real data isn't publicly available.
STATUTORY	A flag derived directly from a codified rule in GFR 2017 or the MPLADS Guidelines.

5 Every Field We Track, In Simple Words

#	Field	What It Means (Simple)	Type
1	Estimated Cost	The amount officially sanctioned for the project.	REAL
2	Actual Expenditure	Worked out from the MP's real yearly fund-release figures, shared across their sanctioned projects (formula in §5A).	DERIVED
3	Payment Amount	Best estimate of how much of the released funds has gone toward this specific project.	DERIVED
4	Physical Progress %	Based on its real official stage (Recommended → Sanctioned → In Progress → Completed).	ESTIMATED
5	Expected Completion Date	Sanction date plus the standard time MPLADS works normally take to complete.	DERIVED
6	Actual Completion Status	A clear, real Completed / Not-Yet-Completed status instead of guessing a date.	REAL
7	Cost Overrun %	How much more is being spent than was originally sanctioned.	DERIVED
8	Delay Days	How many days past the expected finish date an unfinished project already is.	DERIVED
9	Progress Gap	Difference between where a project should be by now and where it actually is.	DERIVED
10	Payment Frequency	How consistently an MP's funds have been released, year after year.	REAL

11	ML Anomaly Score	Isolation Forest looks at all fields above together and gives each project a single outlier score.	REAL AI
12	Agency Risk Score	Graph model scores how structurally unusual an agency's position is across constituencies/MPs.	REAL AI
13	Guideline Violation Flag	Flags statutory ineligible works or a shortfall against the confirmed 15% SC / 7.5% ST allocation quota.	STATUTORY
14	March Rush Index	Flags fiscal year-end sanction clumping using real Finance Ministry expenditure-pacing caps, applied at scheme level as our own calibrated rule.	DERIVED
15	Duplicate Likelihood %	Probabilistic match against Bhuvan/MGNREGA assets in the same block (architecture upgraded in §7).	ESTIMATED
16	Benford Deviation Score NEW	How far an agency's or MP's cost figures deviate from the natural leading-digit distribution — a forensic-accounting fabrication signal (§6E).	DERIVED
17	Completion Probability Curve NEW	Survival-model estimate of the likelihood a project finishes within N months, given agency and sector history (§6F).	DERIVED

5A. The Exact Formula Behind "Derived Expenditure"

A pro-rata allocation: a project that represents a bigger slice of the MP's sanctioned work gets a proportionally bigger slice of whatever funds the MP actually had released that year.

DERIVED EXPENDITURE (Project P, Year Y) = [Sanctioned Cost of P ÷ Total Sanctioned Cost of MP's Active Projects in Y] × MP's Total Funds Released in Y

Worked example (illustrative only): An MP has 3 active projects sanctioned at 80L, 40L, and 30L (total 150L), and had 90L actually released that year. Project A's derived expenditure = $(80 \div 150) \times 90 = 48L$. The same ratio drives Cost Overrun %, Payment Amount, and Progress Gap for every project — one formula, shown to the officer on click.

6 How Our AI & Detection Engines Work — Six Independent Signals

Six signals that don't overlap: a project or agency has to be caught by more than one to be high-priority. That redundancy is what makes the Priority Score defensible, not just a single model's opinion.

6A. Isolation Forest — Individual Project Anomalies

Most fraud-detection systems need thousands of confirmed-fraud examples to learn from. Nobody has that data publicly, so we don't pretend we do. Isolation Forest looks at every project's numbers — cost, progress, delay, similarity to other projects — all at once, and finds the ones that don't fit the normal pattern: the ones easiest to isolate from the crowd.

6B. Graph Theory — Agency-Constituency Network

A lot of real-world fraud is visible in the relationships between projects, not any one project alone. We build a network of agencies and constituencies/MPs, weighted by project count and value, and compute three graph-theory measures per agency: degree centrality (how many disparate constituencies an agency touches), weighted concentration (how much of one MP's funds land with one agency), and betweenness (agencies sitting unusually central across otherwise unconnected MPs). Extreme-tail agencies get their own Agency Risk Score, shown as an interactive node-link graph — rendered at demo-safe scale using the architecture in §7.

6C. Statutory Compliance Engine

Government officers cannot act on statistics alone; they need statutory backing. This engine checks projects against real legal guardrails: the negative list of ineligible works, the 15% SC / 7.5% ST allocation quota, and the year-end expenditure rush check (grounded in Ministry of Finance pacing caps of 33% Jan–Mar / 15% March, applied to MPLADS sanctioning as our own calibrated design choice — which is why that flag is labeled DERIVED, not quoted statutory text). From v5 onward, every one of these rules is versioned and effective-dated per §8, so a guideline revision never breaks a past finding.

6D. Cross-Scheme Block-Level Asset Deduplication

A pervasive drain on public funds is double-invoicing: funding the same rural road or hall under both MPLADS and MGNREGA. We match records down to State + District + Block/GP, then run entity resolution on asset titles within each block. The matching architecture is upgraded in §7 to scale past naive pairwise comparison. Pairs above a tuned similarity threshold in the same Gram Panchayat are flagged with a Duplicate Likelihood warning and shown side-by-side.

MASTERSTROKE · NEW IN V5

6E. Forensic Digit Analysis — Benford's Law

Naturally occurring financial figures follow a predictable leading-digit distribution: $P(d) = \log_{10}(1 + 1/d)$, so "1" leads roughly 30% of the time and "9" leads only about 4.6% of the time. Fabricated or manually padded figures — inflated cost estimates, systematically rounded sanctions — reliably distort this curve. We run a chi-square goodness-of-fit test on the leading digits of sanctioned costs and derived expenditure, grouped by agency and by MP. A significant deviation from the expected Benford curve is a real forensic-accounting red flag, used by tax authorities and auditors worldwide, and it is statistically independent of both the ML anomaly score (multivariate outliers) and the graph signal (relationship structure) — a genuinely separate kind of evidence, shown in the dossier as an expected-vs-actual digit histogram.

Why it matters for the score: this is the one signal a judge is unlikely to have seen in a student project before, it costs almost nothing to compute, and it is fully explainable — no black box, consistent with our honesty framework.

MASTERSTROKE · NEW IN V5

6F. Survival Analysis — "Zombie Project" Early Warning

The problem statement explicitly asks for predictive insights, not just after-the-fact anomaly flags. Instead of a plain trend line, we model project completion as a time-to-event problem using Cox

Proportional Hazards (via the `lifelines` library), conditioned on agency track record, sector, district, and funding pace. The output is a completion-probability curve per project — not a single guess, but a statistically grounded likelihood of finishing within 3, 6, or 12 months. Projects whose hazard curve flattens near zero are surfaced as "zombie project" risk months before a naive delay-days counter would catch them, giving officers real lead time to intervene.

Why it matters for the score: it directly answers the problem statement's "predictive insights" requirement with real statistical modeling instead of linear extrapolation.


7 Engineering for Scale — The Feasibility Architecture

This section exists because "73,000+ records and real-time graph rendering" was named directly as a feasibility risk. Nothing here is decorative — every piece maps to a specific failure mode.

MASTERSTROKE · NEW IN V5

A Bifurcated Compute Architecture — Cold Path + Hot Path

The mistake most teams make is trying to compute everything live, in front of judges. We separate what needs to be pre-computed from what genuinely benefits from being live.



DuckDB is an in-process columnar analytical engine — effectively "SQLite for analytics" — that handles tens of millions of rows in sub-second aggregate queries with no server to manage. Every dashboard click becomes a flat lookup against a pre-materialized cube, not a live pandas operation over 73,000 rows. The hot path never refits a model on stage; it scores one new sample against an already-fitted forest, which is honest about what's real-time and what isn't.

MASTERSTROKE · NEW IN V5

Ego-Graph Rendering — Never Draw the Whole Network

NetworkX is pure Python and was never built for interactive rendering of thousands of nodes — that is almost certainly the actual crash risk in a live graph demo. Two fixes, layered: (1) compute

centrality and betweenness offline with **igraph** or **graph-tool** — C-backed libraries roughly 10–100x faster than NetworkX at the same computation, keeping NetworkX only as the prototyping layer; (2) never render the full national graph. When an officer clicks an agency, render only its 1–2 hop neighborhood — the same "focus and context" pattern used by real graph-intelligence tools like Palantir and Neo4j Bloom. Louvain community detection pre-partitions the graph offline, so a cluster expands on demand instead of rendering thousands of nodes at once. Rendering 12 nodes instead of 5,000 is what actually keeps a live demo from freezing.

MASTERSTROKE · NEW IN V5

Probabilistic Record Linkage at Government Scale

Naive pairwise fuzzy matching, even inside a State+District+Block, degrades badly in dense blocks. There is a strong, verifiable precedent worth building toward: **Splink**, an open-source probabilistic record-linkage library built by the **UK Ministry of Justice** specifically to deduplicate millions of government records at scale, using Fellegi-Sunter probabilistic matching with blocking rules and a DuckDB backend. Adopting the same architecture pattern another government already uses in production for this exact problem is a stronger, more credible claim than a hand-rolled matcher — and it drops directly into the DuckDB serving layer above.

Demo-day fallback: ship the existing block-then-fuzzy-match approach as the guaranteed Core path; present the Splink upgrade as the designed production path. Never let an infrastructure upgrade become a single point of failure on stage.

8 Engineering for Time — The Sustainability Architecture

This section exists because "handling changes in government statutory rules over time" was named directly as a judging weakness. The fix is architectural, not a promise to "update the code later."

MASTERSTROKE · NEW IN V5

Policy-as-Data with Bitemporal Rule Versioning


Every statutory rule — the SC/ST quota percentage, the negative list, the pacing caps — is externalized into a versioned, declarative configuration that sits outside the application code, instead of being hardcoded as an `if` statement. The critical piece: each rule carries an `effective_from / effective_to` date range, so a project sanctioned in 2019 is always checked against the rules that were actually in force in 2019 — not whichever guideline is live today. This is a bitemporal design: it makes the system durable across every future guideline revision with zero code redeploys, and it makes every flag defensible years later in an actual audit.

Rule ID	Clause	Version	Effective Range	Source
SC-ST-QUOTA	15% SC-area / 7.5% ST-area allocation	v2023.1	2023-04-01 → present	MPLADS Guidelines 2023
NEG-LIST-04	Places of worship ineligible	v2023.1	2023-04-01 → present	MPLADS Guidelines 2023, Annexure (pending final citation check)
PACE-CAP	Max 33% Jan–Mar / 15% March release	v2017.1	2017-03-07 → present	Ministry of Finance expenditure-pacing norms

This table is illustrative of the schema, not a claim of completeness — the point being demonstrated is that a rule change becomes one new row, not a code change.

MASTERSTROKE · NEW IN V5

The Guideline Watchdog — Human-Gated, Never Auto-Applied



This keeps the human-in-the-loop principle intact at the *policy* layer, not just the decision layer — a level of rigor that maps directly onto how real regulatory systems are actually maintained, and it feeds the same Rule Provenance data straight into Explain-My-Score: every flag can cite the exact rule ID, version, and source document that produced it.

9 What Happens After a Project Gets Flagged

#	Step	What Happens
1	System flags a project	Based on the Priority Score, combining all six signals from §6.
2	Officer sees "Explain-My-Score"	Plain-language, itemized additive points — e.g. +30 CPWD variance, +25 agency concentration, +20 year-end rush, +15 Benford deviation. Complete transparency, now including exact rule version cited.
3	One-click audit-ready dossier (PDF)	Print-ready inspection memorandum with statutory citations and project evidence pre-filled — built to the level of detail an auditor would expect, not claimed as an official CAG format.
4	Officer requests/conducts field check	District engineer inspects the physical asset on ground with a pre-filled checklist.

5	Officer submits formal recommendation	Logged with verified digital credentials, notes, and photos, creating an immutable audit trail.
6	The existing government system acts	Our system never freezes funds or takes action unilaterally; formal decisions stay with the competent constitutional authority.

10 Who Sees What — Role-Based Dashboards

The problem statement names four specific audiences by name. One underlying dataset, filtered by role — no extra pipeline needed, just permission-scoped views over the same DuckDB serving layer from §7.

Role	What They See	Why It Matters
Member of Parliament	Own constituency's projects, Priority Scores, plain-language flags, SC/ST quota tracker.	A self-check view, not a policing view — helps MPs catch problems before an audit does.
District Authority	Every project physically inside their district across all contributing MPs, interactive map, one-click audit-ready dossier.	For on-ground verification, field-inspection dispatch, agency oversight.
State Nodal Authority	State-wide roll-up — every MP and district, ranked by flag count, agency concentration, unspent funds.	For state-level oversight, cross-district cartel tracking, resource allocation.
Ministry / DIID	National view — cross-state comparison, top flagged agencies and MPs, fund-utilization velocity and trends.	For policy-level monitoring, exactly as the problem statement asks.

11 Why This Approach Wins — The Scorecard

Evaluation Vector	Typical Hackathon Submission	MPLADS-SATARK (v5)
Data Integrity	Quietly uses synthetic or fake data under the hood.	100% real MoSPI (73k+ records) + Bhuvan + CPWD rate schedules.
Data Provenance	Unclear where numbers came from; easily discredited.	Honest REAL / DERIVED / ESTIMATED badges on every data point, including the two new v5 fields.
AI Architecture	Single generic ML model or simple thresholding.	Six independent signals: Isolation Forest, graph theory, statutory engine, deduplication, Benford forensic analysis, survival-model prediction.
Statutory Backing	Ignored, invented, or misquoted financial rules.	GFR 2017 & MPLADS Guidelines 2023, bitemporal-versioned so past flags stay valid across future rule changes.
Explainability	Black-box decimal anomaly score.	Explain-My-Score additive waterfall + explicit pro-rata formula + exact rule citation per flag.

Feasibility at Scale	Untested claims about handling large datasets live.	Named, load-bearing architecture: DuckDB serving layer, pre-fitted hot-path scoring, ego-graph rendering, probabilistic record linkage.
Sustainability	Rules hardcoded; breaks on the next guideline revision.	Policy-as-Data with bitemporal versioning and a human-gated Guideline Watchdog.
Predictive Rigor	A trend line, presented as "AI prediction."	Cox Proportional Hazards survival modeling for genuine completion-probability forecasting.
Field Actionability	Vague "flagged" list with no administrative utility.	One-click, audit-ready printable inspection dossier.
Stakeholder Scope	Single generic "admin" dashboard.	Dedicated dashboards for all four named stakeholders.
Legal Realism	Unrealistically claims AI will "freeze government funds."	Strictly respects constitutional human-in-the-loop governance.

12 The Build Plan — 8-Day Roadmap

Every deliverable is tagged **CORE** — must work live on demo day — or **STRETCH** — valuable, safe to cut or mock if time runs short — so the demo never depends on all seventeen fields working at once.

Timeline	Deliverables
Days 1–2	Ingest MoSPI, Bhuvan/MGNREGA, CPWD datasets CORE . Clean schema, normalize agency names, compute Derived Expenditure CORE . Stand up the DuckDB serving layer CORE — this simplifies days 3–6 rather than adding to them.
Days 3–4	Isolation Forest anomaly score CORE ; NetworkX/igraph agency graph with Louvain pre-clustering CORE ; negative-list & SC/ST quota check as versioned Policy-as-Data rules CORE ; March-rush pacing check CORE ; Benford's Law digit analysis CORE — cheap, high-signal, build it early; block-then-fuzzy-match deduplication CORE ; Splink-based probabilistic linkage upgrade STRETCH .
Days 5–6	4 role-based dashboards CORE ; interactive GIS map view CORE ; Explain-My-Score waterfall modal with rule-citation display CORE ; ego-graph click-to-expand visualizer CORE ; full Cytoscape network view STRETCH ; survival-analysis completion-probability curves STRETCH .
Days 7–8	Recommendation workflow & audit log CORE ; live hot-path injection demo (pre-fitted model scoring a new record) CORE ; end-to-end testing CORE ; audit-ready PDF dossier export STRETCH ; Guideline Watchdog diff-review workflow STRETCH ; demo rehearsal CORE .

If every Stretch item slips, the Core path alone still demonstrates every claim in the §11 scorecard end-to-end — including both new v5 architecture sections.